

Extender-DB – Status Briefing

The coverage milestone is complete: every extender is now mapped to the jobs it serves.

dotfiles-agents · branch feat/extender-db · 2026-07-20 · M1 delivered

THE ESTATE, AT A GLANCE

37 extenders

9 frameworks


1,550 assessments

24 jobs mapped

What this project is

Extender-db is a queryable database of the 37 agent extenders in the marketplace, joined to the mental models used to judge and compose them.


- Two halves joined by assessments: the inventory (what each skill, agent, hook, and external IS) and the doctrine (frameworks: archetypes, taxonomies, authoring specs, the jobs taxonomy, adopted eval methodologies).
- The design bet is "doctrine as data": mental models live as rows with provenance and status, so curation questions – covered, duplicated, missing – are queries, not impressions.
- The repo stays the source of truth. The database is a rebuildable projection.
- It lives as a `_meta/` desk tool on its own branch; promotion to more is an explicit, gated decision that has not been made.




Where it stands

Track I – our own IP – has its foundation complete as of today; Track II – adopted evaluation – is loaded but not yet exercised.

DELIVERED *W0 through M1 – the shared foundation*



NEXT *Track I – our IP (blue) Track II – adopted eval, staged (purple)*



CHARTER PROMOTION GATE

2 of 4 criteria ticked – judged pass ✓, hooks end-to-end ✓; findings→actions and update-path proof remain.

What the database holds now

One thousand five hundred fifty assessment rows now connect the 37-extender inventory to nine frameworks, and every verdict resolves to the exact prompt and response that produced it.


Layer	Contents
Inventory	37 extenders (23 authored skills · 4 agents · 4 hooks · 6 externals), full file contents, packaging, 10-publisher trust registry
Doctrine	9 frameworks, 72 elements – including the 24-job taxonomy and two adopted eval methodologies (SkillOpt, ClosedLoop judges)
Assessments	1,550 rows across mechanical, judged, review, and coverage passes – all provenance-linked
Coverage (new)	24 job_coverage rows: 19 covered · 3 partial · 2 gaps
Relationships (new)	46 rows: 24 complementary · 1 duplicative · 21 directed hand-off links

EVERY VERDICT RESOLVES TO ITS EXACT PROVENANCE



The coverage matrix headline


Of the 24 jobs the estate needs done, 19 are covered by at least one extender whose primary purpose is that job; three are only partially served; two are open gaps.



Deepest coverage: sustain-continuity (3 present + 1 partial), produce-diagram (4 present), integrate-knowledge-system (4 present), manage-backlog and enforce-standards (3 present each).

Every hole has an in-estate answer


The punchline of the coverage work: nothing missing needs to be authored from scratch – every gap and weak spot already has a candidate in the estate.



These five calls are yours – issue #153, each with a default so a bare thumbs-up accepts all.

How M1 was verified (the crew)

The mapping was produced by a layered crew, not a single pass: cheap judges fanned out, blind reviewers re-derived a sample, and the session adjudicated every disagreement.



Full provenance: three eval_runs hold every exact prompt and verbatim response (W7 invariant).
The 662 pre-existing assessment rows survived the live schema change checksum-identical.

What verification actually changed

Blind review is recall, not just error-correction – it added coverage the fan-out missed, and produced boundary rules that become the next rubric.


10 DISPUTED VERDICT CELLS ON THE 8-EXTENDER SAMPLE



- 10 verdict cells disputed on the 8-extender sample; 7 were adjacent-step (the recurring partial boundary, same as W1).
- The reviewers found a real capability the judges missed: readme-value-and-proof's headless-Chromium screenshot capture turned operate-browser-ui from a hard gap into partial coverage.
- Two thin-description externals were reassigned outright – evidence that judging externals by a one-line description is under-determined (now tracked as EDB-23).
- Four boundary rules were logged (for example: consuming the handoff is not sustaining continuity; nudge-only hooks cap at partial). They feed the M4 rubric directly.

The one confirmed redundancy

Only one duplicative pair survived adversarial checking, and the check produced a reusable test for duplication claims.




The doctrine: "produces the same deliverable" is a weak duplication signal. The real test is "same job the SAME WAY" – check for base+override, orchestrator+worker, and hub-routing before believing a duplicative call.

Relationships: delegation & continuity

Beyond coverage, M1 recorded how the extenders hand off to each other – the raw material a future composition planner will reason over.

46 relationship rows: 24 complementary · 1 duplicative · 21 directed hand-off links

→ precedes - -> feeds-into



build loop – scout precedes builder; builder and reviewer feed each other; lead orchestrates




continuity chain – each guard precedes the next, passing the baton forward


Relationships: deliverables, governance & knowledge

The long-game consumer, deferred as EDB-22, is a grounded planner that reads coverage plus directional links plus triggers and assembles an ordered toolkit per goal at request time.

DELIVERABLES



GOVERNANCE



KNOWLEDGE SYSTEM



EDB-22 – the consumer

a grounded planner assembles an ordered toolkit per goal from coverage + links + triggers

Roadmap: two tracks, six milestones

The strategy is to build only what is uniquely ours – the taxonomy, coverage, and curation decisions – and adopt or reuse everything else.

TRACK I – our IP



TRACK II – adopted eval (staged)



Execution principle (charter decisions 6-8): methodologies are adopted (SkillOpt, ClosedLoop judges), the eval harness is reused – M5 runs through Henry's existing meta-harness on cheap models. Nothing bespoke gets built.

Decisions made – and the plan forward

Henry decided the batch on July twentieth: items A through E accepted as recommended, and F reversed – the saved gaps view will be built.

A	The five roster promotions from slide 6.	approved
B	Drop / fold the duplicative coleam00 excalidraw skill.	approved – drop
C	Accept vendor-only coverage for improve-code via anthropic's code-simplifier.	accepted
D	Approve the obsidian-cli reference fix next session – satisfies BOTH remaining promotion-gate criteria.	approved
E	The repo/path pointer to the meta-harness – M5 cannot start without it, and nothing gets built in its place by design.	STILL NEEDED – blocks M5
F	Optional: a saved gaps view in the database.	REVERSED – will build

All remaining development is now tracked as epic #154 with 13 sub-issues (#155-#167) – promotions, coleam00 drop, gate path, coverage_gaps view, M2-M6.

Health and invariants

Everything the project asserts about itself was verified by gates this session, and the operating knowledge is now written down.

ALL DEFINITION-OF-DONE GATES GREEN

✓ idempotent loaders

✓ ingest checksum-verified

✓ zero unlinked assessments

✓ make ci – 120 tests


✓ data.db WAL-
checkpointed

- Standing invariants: session-only database credentials (judge agents are structurally read-only); every verdict provenance-linked; generated artifacts have regen scripts; schema changes merge-by-name.
- New this session: PROCEDURES.md – the runbook with script run order, the evaluated-pass pattern, gates, and the commit discipline – so any cold session can operate the project.

Status in one line

Foundation complete and verified; decisions made; all remaining development tracked on epic #154 – the evaluation track starts once the meta-harness pointer lands.

THE PATH FROM HERE



Branch feat/extender-db pushed through fa84ecd · working tree clean · server restartable via serve.sh

Epic: #154 · Runbook: PROCEDURES.md · Matrix: coverage-matrix.md